GLOBAL TRUST CHALLENGE
Building Trust in the Age of AI

Global Trust Challenge (2025)
Phase 1 - Proposal

NewsTrace: A Provenance and Transparency System for AI-Assisted Journalism

woVxPrRV

Entry details

Briefly describe your proposed approach *

> Artificial intelligence is increasingly being used to write, edit, and summarize news articles. While this can make journalism faster and more efficient, it creates a serious problem: readers have no way of knowing how much of what they are reading was produced by an AI system versus a human journalist. This opacity is eroding public trust in digital media at a time when trust is already fragile.
>
> NewsTrace proposes a three-part transparency infrastructure to address this problem. The first component is a Provenance Recorder — a lightweight tool that operates in the background during the writing process, logging which parts of an article were typed by a human and which were generated or suggested by an AI tool. The second component is an Editorial Gatekeeper — a configurable rules engine that allows newsrooms to enforce their own AI use policies before an article is published, ensuring that AI-generated content is always reviewed by a human editor. The third component is a Reader Transparency Label — a public-facing panel embedded in published articles that shows readers a clear, simple breakdown of AI involvement, similar to the way a nutrition label shows the ingredients in food.
>
> Together, these three components create a complete trust pipeline: from creation to publication to public disclosure. The system is designed to be open-source, lightweight, and adoptable by newsrooms of any size anywhere in the world.
>
> This proposal focuses on the design and conceptual framework of NewsTrace, with a development roadmap for future implementation. The goal is to establish a reference model for AI transparency in journalism that policymakers, technologists, and news organizations can build upon.

Primary domain of your proposal *	Integrated (Policy+Tech)

Proposed Policy Approach *

The policy challenge I am addressing is the absence of enforceable AI disclosure standards in digital journalism.

Governments around the world are beginning to recognize that AI-generated content poses a serious threat to information integrity. The EU Artificial Intelligence Act (2024) requires that AI-generated content be disclosed to the public. Several US states have introduced similar legislation. However, these laws share a critical weakness: they mandate disclosure without providing any technical standard for how that disclosure should work in practice. A law that says "you must label AI content" is meaningless if there is no reliable, verifiable way to do it.

The result is a patchwork of inconsistent, unverifiable disclosures. Some newsrooms add a single line at the bottom of an article saying "AI tools were used in the production of this content." This tells the reader almost nothing. It cannot be verified. It can be removed. It gives no detail about the degree or nature of AI involvement. Readers are left guessing — and trust erodes.

The policy concept underpinning NewsTrace is provenance-as-disclosure.

Instead of treating AI disclosure as a label added after publication, NewsTrace proposes that disclosure should be built into the content creation process itself — recorded at the moment of writing, attached permanently to the article, and made publicly verifiable. This transforms disclosure from a voluntary afterthought into a structural, tamper-resistant feature of every published article.

This approach draws on models from other industries. Food packaging laws do not just say "disclose ingredients" — they specify exactly what must be disclosed, in what format, and in a standardized way that consumers can understand and compare. Financial reporting laws require auditable records, not just summary statements. News Trace proposes that AI disclosure in journalism should meet the same standard: structured, verifiable, and permanent.

The specific policy innovation is a three-layer disclosure framework:

First, newsrooms must maintain an internal provenance log — a record of which parts of every article were human-written versus AI-generated. Second, an editorial review requirement must be met before publication — ensuring a human is accountable for every published piece. Third, a standardized public transparency label must be attached to every article — giving readers consistent, comparable information across all news outlets.

This framework is designed to be technology-neutral and globally adoptable. It does not prescribe which AI tools newsrooms may use. It simply requires that whatever is used is disclosed honestly, verifiably, and in a format the public can understand.

At a time when public trust in journalism is at a historic low, and AI is accelerating the production of content that readers cannot distinguish from human writing, this policy approach offers a concrete, enforceable, and scalable path toward rebuilding that trust.

Evidence Base and Relevance *

NewsTrace builds directly on three existing frameworks while addressing a gap that none of them currently fill.

1. C2PA — Coalition for Content Provenance and Authenticity C2PA is an open industry standard developed by Adobe, Microsoft, BBC, and others to cryptographically attach provenance metadata to digital media. It is already being used to label AI-generated images and videos. However, C2PA does not yet address text-based journalism — a significant gap given that most news is still consumed as written articles. NewsTrace applies the same provenance logic C2PA uses for images, but specifically designed for the text writing workflow.

2. W3C PROV Standard The World Wide Web Consortium's PROV specification provides a formal model for recording who produced what, when, and how on the web. NewsTrace's Provenance Recorder is designed around this model

> — every contribution to an article is recorded as an entity (the text), an activity (writing or AI generation), and an agent (human or AI tool). This gives the system a rigorous, internationally recognized technical foundation.
>
> 3. EU Artificial Intelligence Act (2024) The EU AI Act legally requires that AI-generated content be disclosed to the public. However, the law provides no technical standard for how newsrooms should implement this disclosure. A 2024 Reuters Institute report found that over 60% of news organizations are already using AI in their content workflow, yet consistent disclosure practices remain absent. NewsTrace proposes exactly the missing technical infrastructure the law demands but does not specify.
>
> The Gap NewsTrace Fills Existing voluntary disclosure practices — a line at the bottom of an article saying "AI-assisted" — are unverifiable, inconsistent, and easily removed. No current framework makes AI provenance a permanent, tamper-resistant, and reader-friendly part of the article itself. NewsTrace is designed to fill that gap by combining the cryptographic rigour of C2PA, the formal model of W3C PROV, and the public disclosure intent of the EU AI Act into a single, open-source, journalism-specific system.

Implementation Considerations

Target jurisdiction or sectors *

> This proposal targets the global digital journalism sector, with initial focus on English-language online newsrooms. It is designed to align with the EU Artificial Intelligence Act and emerging US state-level AI disclosure laws. The open-source approach ensures it can be adopted across jurisdictions, including lower-income countries with limited regulatory infrastructure.

Stakeholders needed *

> Three groups are essential. First, newsrooms and journalists — as the primary users of the system, their workflow needs must shape the design. Second, technology developers and standards bodies such as C2PA and W3C — to ensure NewsTrace integrates with existing infrastructure. Third, policymakers and regulators — to validate that the transparency label meets legal disclosure requirements and to encourage adoption across the industry.

Governance implications *

> NewsTrace raises three governance considerations. First, who owns and audits the provenance logs — newsrooms, independent bodies, or regulators? Second, how are disputes resolved when a log is contested? Third, how is the system governed as an open-source standard to prevent fragmentation across different implementations. These questions suggest the need for a neutral, multi-stakeholder governance body — similar to how C2PA is governed today.

If applicable, does your approach experiment with novel governance models?

> NewsTrace proposes a distributed accountability model where governance of AI transparency is shared across three layers — the newsroom (through configurable editorial policies), the technology (through tamper-evident provenance logs), and the public (through the reader-facing label). Rather than relying on a single regulator to enforce disclosure, the system embeds accountability directly into the journalism workflow itself.

Technology approach description *

> NewsTrace proposes a three-component software system designed to make AI involvement in journalism transparent, verifiable, and publicly accessible.
>
> Component 1 — Provenance Recorder A lightweight browser extension that runs silently in the background while a journalist writes. It monitors contributions to the article in real time, tagging each piece of text with a source label: HUMAN if typed by the journalist, or AI-GENERATED / AI-ASSISTED if produced or suggested by an AI tool. Each tagged contribution is timestamped and cryptographically hashed using SHA-256, creating an append-only log that cannot be altered after the fact. This log forms the verified record of how the article was produced.

> Component 2 — Editorial Gatekeeper A rules engine integrated into the newsroom's publishing workflow. Before an article can be published, the system checks whether the newsroom's own AI use policies have been followed — for example, confirming that all AI-drafted sections were reviewed by a human editor. Policies are defined in a simple configuration file, making the system adaptable to any newsroom's standards without requiring custom development.
>
> Component 3 — Reader Transparency Label An embeddable widget attached to every published article. Readers can click it to see a plain-language breakdown of the article's provenance — what percentage was human-written, what percentage was AI-assisted, which AI tools were used, and whether editorial review took place. The label links to the full provenance log for readers who want deeper verification.
>
> The proposed technology stack is deliberately accessible: JavaScript and TypeScript for the browser extension, React for the reader-facing label, and Node.js for the backend provenance log API. The entire system would be released as open-source software, ensuring any newsroom in the world can adopt it at no cost.
>
> This approach builds on the C2PA standard for media provenance and the W3C PROV data model, extending their principles specifically to text-based journalism for the first time.

Key Features & Capabilities

What technical features distinguish your approach?

How do these features advance trust?

Trustworthiness mechanisms (e.g., provenance, explainability, transparency) *

> NewsTrace embeds trust at three levels. At the creation level, the Provenance Recorder cryptographically hashes every contribution using SHA-256, making the authorship log tamper-resistant and independently verifiable. At the editorial level, the Gatekeeper enforces human review before publication, ensuring AI content is never published without accountability. At the reader level, the Transparency Label provides plain-language disclosure of AI involvement directly on the article. Together these mechanisms make trust structural — built into the workflow itself — rather than a voluntary afterthought that can be skipped or removed.

Adaptability across cultures or geographies *

> NewsTrace is designed as open-source software with no licensing cost, making it accessible to newsrooms in any country regardless of budget. The Editorial Gatekeeper uses a simple configuration file, allowing newsrooms to define AI policies that reflect their local editorial standards and cultural context. The Reader Transparency Label supports internationalization, meaning the interface can be translated into any language. No centralized authority controls the system — any newsroom can deploy and adapt it independently.

Feedback loops or user safeguards *

> Journalists can review their own provenance logs at any time, giving them visibility into how their work is being recorded. Readers can flag articles where the transparency label appears inconsistent with the content, creating a public accountability mechanism. Newsrooms receive aggregate data on AI usage patterns across their publication, enabling editorial teams to refine their AI policies over time. All user study data collected during development is anonymized, and participation is strictly voluntary.

Technology Maturity * Idea stage

How does your tech promote digital trust and content integrity? *

> NewsTrace promotes digital trust by making the invisible visible. Today, AI involvement in journalism is largely hidden — readers have no reliable way to know whether an article was carefully reported by a human or rapidly generated by an AI system. This uncertainty undermines trust in all news, even the responsibly produced kind.
>
> NewsTrace addresses this directly by attaching a verified, tamper-resistant provenance record to every article. This record is not a voluntary disclosure statement that can be worded vaguely or removed — it is a cryptographic log generated during the writing process itself, making it independently verifiable.
>
> The Reader Transparency Label then translates this technical record into a format any reader can understand, regardless of their technical background. By making AI involvement concrete, specific, and checkable, NewsTrace shifts the foundation of trust from asking readers to believe newsrooms, to giving them the evidence to verify for themselves.

Complementary aspect(s) of the plan *

> NewsTrace complements existing policy and technical frameworks rather than replacing them. On the policy side, it provides the technical implementation layer that the EU AI Act and emerging national disclosure laws require but do not specify — turning legal obligations into practical, auditable actions.
>
> On the technical side, it extends the C2PA provenance standard — currently focused on images and video — into text journalism, filling a documented gap in the existing standard.
>
> At the industry level, NewsTrace complements newsroom editorial guidelines by giving them a digital enforcement mechanism. A newsroom's AI policy is only as strong as its ability to verify compliance — the Editorial Gatekeeper provides exactly that.
>
> Finally, NewsTrace complements media literacy efforts. Helping readers understand AI involvement in the news they consume is itself an educational act — making the transparency label a tool for building long-term public understanding, not just one-time disclosure.

Problem Definition *

> Artificial intelligence is now widely used in newsrooms to draft articles, generate headlines, summaries reports, and produce content at scale. Yet readers have no reliable way of knowing how much of what they read was written by a human versus generated by an AI system. This invisible use of AI is creating a crisis of trust in digital journalism at a moment when media credibility is already fragile.
>
> Three specific challenges make this urgent. First, there is no technical standard for recording AI involvement in text-based journalism at the point of creation. Disclosure, where it happens at all, is voluntary, inconsistently worded, and unverifiable. Second, newsrooms experimenting with AI tools have no shared infrastructure for enforcing their own editorial policies around AI use — meaning internal guidelines exist on paper but cannot be systematically applied. Third, readers have no consistent, trustworthy signal to help them contextualize what they are reading — leading to blanket distrust of all digital news, including responsibly produced journalism.
>
> This issue is underserved for a structural reason: the organizations best positioned to build transparency tools — technology platforms and large media companies — have commercial incentives to keep AI involvement opaque. Independent, open-source infrastructure that serves the public interest rather than platform interests does not yet exist for text journalism.
>
> The urgency is compounded by regulation moving faster than implementation. The EU AI Act now legally requires AI content disclosure, and similar laws are emerging globally. But no technical standard exists to help newsrooms comply in a consistent, verifiable way. This gap between legal obligation and technical reality is exactly what this proposal addresses.

Proposed Approach *

NewsTrace proposes a three-component transparency infrastructure for AI-assisted journalism, designed to make AI involvement visible, verifiable, and publicly understandable at every stage of the content lifecycle — from creation through publication to public consumption.

Component 1 — Provenance Recorder The Provenance Recorder operates silently during the writing process. As a journalist works, it logs every contribution to the article in real time, tagging each piece of text as HUMAN, AI-GENERATED, or AI-ASSISTED. Each entry is timestamped and cryptographically hashed using SHA-256, creating an append-only provenance log that cannot be altered after the fact. This log is the foundational trust mechanism — it transforms AI disclosure from a voluntary statement into a verifiable technical record generated at the moment of creation.

Component 2 — Editorial Gatekeeper The Editorial Gatekeeper is a rules engine embedded in the publishing workflow. Before an article can be published, it checks whether the newsroom's AI use policies have been followed — for example, confirming that AI-drafted sections were reviewed by a human editor, or that factual claims generated by AI were independently verified. Policies are defined in a simple configuration file, making the system adaptable to any newsroom's editorial standards without custom development. This mechanism introduces accountability inside the organization before anything reaches the public.

Component 3 — Reader Transparency Label The Transparency Label is a public-facing widget embedded in every published article. Readers can click it to see a plain-language breakdown of how the article was produced — what percentage was human-written, what percentage was AI-assisted, which AI tools were involved, and whether human editorial review took place. It links to the full provenance log for readers who want to verify the record independently. The label is designed to be understood by any reader regardless of technical background, drawing on the familiar model of a nutritional information panel.

Trust Mechanisms NewsTrace builds trust through four mechanisms: provenance — a tamper-resistant record of who produced what; explain ability — plain-language disclosure that any reader can understand; accountability — mandatory editorial checks before publication; and auditability — a publicly accessible log that regulators, researchers, and readers can independently verify.

Intended Users and Beneficiaries The primary users are journalists and editors, who gain a structured way to document and enforce responsible AI use. The primary beneficiaries are readers — particularly those in high-information environments where distinguishing reliable from unreliable news is critical. Secondary beneficiaries include regulators seeking technical tools to implement AI disclosure law, independent researchers studying AI's role in media, and smaller newsrooms in lower-income countries that currently lack the resources to develop transparency tools independently. As open-source software, NewsTrace is accessible to all of these groups at no cost.

Contribution to Information Integrity and Trust

How does your proposal strengthen information integrity and build trust in digital content and communication? *

NewsTrace strengthens information integrity by addressing the root cause of AI-related distrust in journalism — not the existence of AI in newsrooms, which is already widespread and often beneficial, but the absence of any reliable mechanism for disclosing it.

The system counters disinformation risk in two ways. First, by creating a tamper-resistant provenance log at the point of creation, it makes it significantly harder for bad actors to publish entirely AI-fabricated content without detection. Second, by normalizing transparency as a structural feature of the publishing workflow, it raises the baseline standard across the industry — making opaque AI use increasingly anomalous and therefore more visible.

NewsTrace enhances transparency by moving disclosure from voluntary and vague to structural and verifiable. A statement at the bottom of an article saying "AI-assisted" is unverifiable and easily ignored. A cryptographically signed provenance log attached to the article and surfaced through a reader-facing label is neither.

> It reinforces accountability by embedding editorial responsibility into the publishing process itself. The Gatekeeper component means a newsroom's AI policy is not just a document — it is an enforced checkpoint that must be passed before publication.
>
> It fosters user confidence by giving readers something they currently lack entirely — a consistent, trustworthy signal about the content they are consuming. Research consistently shows that transparency, even about imperfect processes, increases rather than decreases trust. Readers do not need journalism to be AI-free. They need it to be honest.
>
> Collectively, these mechanisms shift the foundation of trust in digital journalism from reputation alone to verifiable evidence.

Which aspects of information integrity and trust does your proposal most directly address? *

- ✓ Content authenticity / provenance
- ✓ Platform governance or transparency
- ✓ Trust signals or user literacy
- ✓ Public communication or media integrity
- ✓ Safety of AI-generated content
- ✓ Cross-border policy coherence
- ✓ Data access / auditability

Context and Applicability *

> NewsTrace is designed for immediate applicability in the digital journalism sector, with initial testing most feasible in English-language online newsrooms — particularly student newspapers and independent digital outlets, which are more agile in adopting new tools than large legacy media organizations.
>
> In terms of jurisdiction, the system is directly relevant wherever AI disclosure legislation is active or emerging. The EU is the most immediate context, given the AI Act's disclosure requirements now in force. The United Kingdom, Canada, Australia, and several US states are developing similar frameworks, all of which NewsTrace could support technically.
>
> Beyond journalism, the core infrastructure — provenance logging, policy enforcement, and public disclosure labelling — is adaptable to other sectors where AI-generated content raises trust concerns, including academic publishing, legal documentation, and public sector communications.
>
> Geographically, the open-source design is the key adaptability mechanism. Because there is no licensing cost and no centralized controlling authority, any newsroom anywhere in the world can deploy the system. The Editorial Gatekeeper's policy configuration is language-agnostic, and the Reader Transparency Label supports internationalization, allowing the interface to be translated and culturally adapted for any market.
>
> The system is particularly relevant for newsrooms in the Global South, where regulatory infrastructure is developing and where locally built transparency tools are largely absent. By providing a ready-made, freely available foundation, NewsTrace reduces the barrier to responsible AI adoption for resource-constrained newsrooms that would otherwise have no practical path to compliance.

Trustworthiness-by-Design *

> NewsTrace is built around four trustworthiness-by-design principles from the ground up.
>
> Auditability is the foundational principle. The Provenance Recorder generates a cryptographically hashed, append-only log at the moment of writing. This log cannot be altered retroactively, and it is publicly accessible — meaning any reader, researcher, or regulator can independently verify the provenance record attached to any article. Trust is not requested; it is demonstrable.

Explain ability is built into the Reader Transparency Label. The provenance data is translated into plain language that any reader can understand without technical knowledge. The design draws on established communication models — particularly nutritional labelling — to make complex information immediately interpretable.

User control is embedded in the Editorial Gatekeeper. Newsrooms define their own AI use policies through a simple configuration file, retaining full editorial sovereignty over what standards they enforce. The system provides the infrastructure; editorial judgment remains with humans.

Interoperability is ensured by grounding the system in established open standards — specifically the C2PA provenance framework and the W3C PROV data model. This means NewsTrace does not create a proprietary silo. Its provenance logs can be read and verified by any system that supports these standards, and it can integrate with existing content management systems without requiring newsrooms to replace their current infrastructure.

Together these principles ensure that trustworthiness is not a feature added to NewsTrace — it is the architecture.

Governance and Oversight Innovation *

NewsTrace addresses a specific and well-documented governance gap: AI disclosure laws exist but technical implementation standards for text journalism do not. The proposal contributes to governance in three ways.

First, it offers a reference implementation that regulators can point to when defining technical compliance standards for AI disclosure. Rather than leaving newsrooms to interpret vague legal obligations independently, NewsTrace demonstrates what verifiable compliance could look like in practice.

Second, the open-source model itself is a governance innovation. By making the infrastructure freely available and community-maintainable, it creates a public-interest alternative to proprietary disclosure systems controlled by large platforms — reducing the risk of disclosure standards being shaped by commercial incentives.

Third, the Editorial Gatekeeper introduces a soft-law mechanism at the organizational level. Newsrooms can encode their own editorial AI policies into the system, creating enforceable internal standards that complement but do not depend on external regulation. This is particularly valuable in jurisdictions where formal AI regulation is still developing — it allows responsible newsrooms to demonstrate good practice independently of whether a law requires it.

Longer term, the provenance logs generated by NewsTrace create a dataset that researchers and regulators can use to understand how AI is actually being used in journalism at scale — providing an empirical foundation for evidence-based policy that currently does not exist.

Prototype Vision *

A working prototype of NewsTrace developed over 4–6 months would focus on the first and most foundational component — the Provenance Recorder — with a basic version of the Reader Transparency Label as the visible output.

Concretely, the prototype would consist of a browser extension that integrates with Google Docs, monitoring contributions in real time and tagging text as HUMAN or AI-GENERATED. Every tagged contribution would be logged with a timestamp and a SHA-256 hash, stored in a lightweight Node.js backend. A simple web interface would allow the journalist to review their provenance log at any time. The Reader Transparency Label would then read from this log and display a visual breakdown — percentage human-written, percentage AI-assisted — as an embeddable badge on a test article page.

The outputs produced would be: a working browser extension, a provenance log API, a basic transparency label widget, and a short technical specification document describing the architecture for the Gatekeeper component as future work.

Testing would involve two rounds. The first round would be a technical validation — running the extension across 20 to 30 test articles written with varying degrees of AI assistance, verifying that the log correctly captures and hashes all contributions without data loss. The second round would be a usability test with 8 to 10 journalism or communications students, asking them to write using the extension and evaluating whether the provenance log and label accurately reflect how the article was produced.

Potential for Piloting *

If the prototype performs well in testing, the natural pilot setting is a student newspaper. Most universities publish a student-run online newspaper that operates on real editorial deadlines, uses digital writing tools, and is staffed by people willing to participate in research. This makes it an ideal low-risk environment that nonetheless reflects genuine newsroom conditions.

A proposed pilot plan would run over six to eight weeks. In the first two weeks, the NewsTrace extension and label would be deployed across the newspaper's editorial team, with a one-hour onboarding session explaining how the system works. For the following four weeks, all articles produced would run through the Provenance Recorder, and the transparency label would be displayed on the published versions. In the final two weeks, structured feedback would be collected from three groups — journalists using the extension, editors using the Gatekeeper dashboard, and readers responding to the transparency label through a short survey.

Success metrics for the pilot would include: zero data integrity failures in the provenance log, journalist-reported friction score below three on a five-point scale, and at least 70 percent of reader survey respondents reporting that the transparency label increased their understanding of how the article was produced.

Beyond the university setting, the pilot model is directly replicable at community news outlets, fact-checking organizations, and independent digital publications — any organization that produces text content with AI involvement and wants a practical transparency mechanism.

Resources & Partnerships Needed (non-binding): *

- ✓ Expert mentorship
- ✓ Technical partners
- ✓ Policy advisors
- ✓ Testing environment
- ✓ Check if you're open to being matched with other teams for complementary expertise

1. Affected Communities: *

NewsTrace affects four communities, each in distinct ways.

Readers of digital news are the primary beneficiaries. They currently consume journalism with no reliable way to understand AI's role in producing it. NewsTrace gives them a consistent, verifiable signal — the Transparency Label — that restores an element of informed choice to news consumption. This benefit is universal but is most significant for communities that rely heavily on digital news as their primary information source, including younger audiences and people in regions where print journalism has declined.

Journalists and editors are both users and affected parties. They benefit from having a structured, defensible record of their editorial process — particularly valuable as public scrutiny of AI use in newsrooms increases. However, they are also potentially affected by misuse of the audit log as a surveillance tool. To mitigate this, the design process must involve journalists directly. Structured interviews with working journalists and editors will be conducted during Phase 1 to ensure the system serves their needs and respects their professional autonomy.

Smaller and independent newsrooms, particularly in lower-income countries, benefit from having access to transparency infrastructure they could not afford to build independently. The open-source model is a direct response to their exclusion from tools that currently exist only inside large, well-resourced media organizations.

> Policymakers and regulators benefit from a concrete reference implementation they can use when defining technical compliance standards for AI disclosure legislation.
>
> Community involvement is built into the development process. Journalists participate in Phase 1 interviews and Phase 3 usability testing. Reader responses are collected during the pilot. No design decisions about the transparency label or gatekeeper policies will be finalized without input from the people who will use them.

Ethical and Governance Considerations: *

> Risk 1 — Gaming the system A dishonest actor could attempt to manipulate the provenance log by falsely labelling AI-generated content as human-written. Mitigation: The SHA-256 cryptographic hashing of each contribution means any alteration after writing is detectable. Independent auditors or regulators can verify logs without requiring access to the newsroom's internal systems.
>
> Risk 2 — Journalist surveillance The audit log could be misused by newsroom management to monitor individual journalists' productivity or writing patterns rather than for transparency purposes. Mitigation: Privacy-by-design principles are applied from the start. The log records content provenance, not individual performance metrics. Access controls are strictly defined — provenance data is visible to the public at the article level, but granular writing-session data is accessible only to the journalist themselves.
>
> Risk 3 — Reader misinterpretation A reader seeing a high AI-involvement percentage may incorrectly conclude the article is untrustworthy, even if AI was used only for minor editing. Mitigation: The Transparency Label is designed with plain-language contextual explanations of what different types of AI involvement mean, informed by user testing before public deployment.
>
> Risk 4 — Bias in AI detection Automated detection of AI-generated text is imperfect and can produce false positives. Mitigation: NewsTrace does not rely on AI detection algorithms. Provenance is logged at the point of creation — the system records what actually happened, not what it guesses happened retrospectively.
>
> Ongoing Governance During the pilot phase, a small advisory group consisting of a journalism ethics expert, a technologist, and a reader representative will review the system's outputs monthly. Findings will be published openly. Any community member who identifies a discrepancy between a published label and the actual article content will have a clear mechanism to flag it for review.

Team Name *	Apex Trace
Country *	Sri Lanka
Team Lead (Name) *	Chamika Udayanga
Team Lead (Affiliation) *	Faculty of Computing, Sabaragamuwa University of Sri Lanka.
Email Address *	c.udayanga2k04@gmail.com
Total Number of Team Members *	1

Team Description *

> This project is an individual submission by a first-year Computer Science undergraduate. As the sole team member, I am responsible for the full scope of the proposal — problem definition, system design, technical architecture, and development roadmap. I am open to being matched with complementary teammates or mentors through the Global Trust Challenge if the opportunity arises.

Team Composition *	✓ Technical ✓ Youth-led
How did you hear about the challenge?	GTC Partner
I confirm that this submission is original and does not infringe on others’ rights. *	✓
I acknowledge that funding or sponsorship is not guaranteed, even if my proposal is selected, and will depend on external support. *	✓
I understand that if selected, I agree to be bound by the terms and conditions of the global challenge *(will be provided to all that are selected for the prototype phase)* *	✓
I agree to be contacted for future updates and opportunities. *	✓

Log in to app.globalchallenge.ai to see complete entry attachments.

PDF

NewsTrace_Atta... 5 KiB

PDF

NewsTrace_Atta... 151 KiB